

# CS 6476 PS4

Name

GT Email

GT ID

**Part 1: Standard Scaler: Why did we use StandardScaler instead of looping over all the dataset twice for mean and standard deviation? Why a simple loop will not be a good choice in a deployed production grade ML system?**

<text answer here>

**Part 1: Why do we normalize our data (0 mean, unit standard deviation)?**

<text answer here>

**Part 3: Loss function. Why did we need a loss function?**

<text answer here>

**Part 3: [2 pt ]** KL-divergence is defined as

$$D_{KL}(P||Q) = \mathbb{E}_{X \sim P} \left[ \log \frac{P(x)}{Q(x)} \right]$$

Cross-entropy is defined as

$$H(P, Q) = -\mathbb{E}_{X \sim P} \log Q(x)$$

How would you write cross-entropy in terms of KL-divergence

<text answer here>

## Part 5: Training SimpleNet

<Loss plot here>

<Accuracy plot here>

Final training accuracy value:

Final validation accuracy value:

Part 6: Simple Segmentation Net

Class Index	Class name	Simple Segmentation Net Class IoU
0	Building	
1	Tree	
2	Sky	
3	Car	
4	SignSymbol	
5	Road	
6	Pedestrian	
7	Fence	
8	Column_Pole	
9	Sidewalk	
10	Bicyclist	

Validation mIoU:

Number of Epochs: \_\_\_\_\_

## Part 6: Simple Segmentation Net

Paste a figure of the generated semantic segmentation from Colab. It should be a 2x3 grid, with ground truth on the top row, and your predictions on the bottom row.

## Part 6: Simple Segmentation Net

Which classes have the lowest mIoU? Why might they be the most difficult? Provide an example RGB image from Camvid that illustrates your point.

**Conclusion: briefly discuss what you have learned from this project.**

<Text solution here>



# Theory: Neural Nets

1. Given are the layers of a neural network.  
Calculate the total number of learnable parameters for each layer and in total.
  - a) Conv 1: <answer here>
  - b) Conv 2: <answer here>
  - c) FC1: <answer here>
  - d) Total (sum all 3): <answer here>


```
(conv1): Conv2d(3, 64, kernel_size=(5, 5), stride=(1, 1))
(maxpool1): MaxPool2d(kernel_size=2, stride=2, padding=0)
(conv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1))
(maxpool2): MaxPool2d(kernel_size=2, stride=2, padding=0)
(fc1): Linear(in_features=3200, out_features=10, bias=True)
(softmax): Softmax(dim=1)
```

# Theory: Neural Nets (cont.)

Here is the computational graph for the function  $f(x, y, z) = x * e^{y+z}$ . We want to compute the partial derivative of this function with respect to each input. Here are the three intermediate functions

$$A = y + z \quad B = e^A \quad C = x * B$$

From these we build the feed forward computation graph, with inputs  $n, m$ . The top input is  $n$  and the bottom is  $m$ .



a. What is the correct expression for  $\frac{\partial C}{\partial x}$  ?

Type equation here.

b. What is the correct expression for  $\frac{\partial C}{\partial B}$  ?

Type equation here.

c. What is the correct expression for  $\frac{\partial B}{\partial A}$  ?

Type equation here.

d. What is the correct expression for  $\frac{\partial A}{\partial y}$  ?

Type equation here.

e. What is the correct expression for  $\frac{\partial C}{\partial y}$  ?

Type equation here.

### **EC1.1: Screenshot of your `get_data_augmentation_transforms()`**

<Screenshot here if attempted; do not delete the slide if not attempted>

## EC1: Training to solve overfitting

<Loss plot here>

<Accuracy plot here>

Final training accuracy value:

Final validation accuracy value:

EC2: PSPNet

Class Index	Class name	PSPNet Class IoU
0	Building	
1	Tree	
2	Sky	
3	Car	
4	SignSymbol	
5	Road	
6	Pedestrian	
7	Fence	
8	Column_Pole	
9	Sidewalk	
10	Bicyclist	

Validation mIoU:

Number of Epochs: \_\_\_\_\_

## EC2: PSPNet

Paste a figure of the generated semantic segmentation from Colab. It should be a 2x3 grid, with ground truth on the top row, and your predictions on the bottom row.